

Advance DevOps Case Study (Topic 9)

Aim: Static Analysis Integration with Terraform.

- **Concepts Used:** Jenkins, SonarQube, Terraform.
- **Problem Statement:** "Use Terraform to set up the infrastructure for a Jenkins server and a SonarQube instance. Then, configure a Jenkins pipeline to perform a static analysis of a Python application using SonarQube."
- **Tasks:**
 - Write a Terraform script to deploy Jenkins and SonarQube on AWS.
 - Set up a Jenkins pipeline to analyze a Python codebase using SonarQube.
 - Trigger the pipeline and review the analysis results in SonarQube.

Introduction:

- **Case Study Overview:**

This case study illustrates how to integrate static code analysis into a DevOps pipeline using Terraform, Jenkins, and SonarQube. By utilizing Terraform, the deployment of Jenkins and SonarQube servers on AWS is automated, ensuring a consistent and scalable infrastructure.

Jenkins serves as the CI/CD server, automating tasks like building and testing code, while SonarQube analyzes the code for bugs, vulnerabilities, and code smells. The setup provides immediate feedback to developers, enabling early detection of issues.

With Terraform's Infrastructure-as-Code approach, the entire process is streamlined, reducing manual setup and ensuring a reliable environment. This case study will demonstrate how these tools work together, from deploying infrastructure to configuring the Jenkins pipeline and reviewing SonarQube analysis results, promoting better code quality and efficient development practices.

- **Key Feature and Application:**

The unique feature of this integration is the automated deployment and configuration of the CI/CD pipeline and code analysis tools using Terraform, making the setup reproducible and

scalable. This infrastructure-as-code (IaC) approach ensures that every part of the setup—from servers to configuration—can be managed and version-controlled, leading to a more reliable and consistent environment.

Practical Use:

- **Enhanced Code Quality:** Developers can receive immediate feedback on code issues, such as bugs, code smells, and vulnerabilities, during the CI/CD process.
- **Automated Deployment:** Utilizing Terraform allows for quick and consistent deployment of the Jenkins and SonarQube infrastructure, reducing the manual effort required.
- **Scalability and Efficiency:** Since all components are managed as code, the infrastructure can be easily scaled or replicated, making it suitable for dynamic environments.

Prerequisites:

1. An AWS account with an IAM user set up on it.
2. Terraform downloaded on your system.

Step-by-Step Explanation:

Step 1: Create a folder for your project and create the following files inside of it:-

a. main.tf:-

```
resource "aws_instance" "jenkins" {
  ami          = var.jenkins_ami_id
  instance_type = var.instance_type
  key_name     = var.key_pair_name
  security_groups = [aws_security_group.jenkins_sg.name]

  tags = {
    Name = "Jenkins Server"
  }
}
```

```
resource "aws_instance" "sonarqube" {
  ami          = var.sonarqube_ami_id
  instance_type = var.instance_type
  key_name      = var.key_pair_name
  security_groups = [aws_security_group.sonarqube_sg.name]

  tags = {
    Name = "SonarQube Server"
  }
}
```

b. outputs.tf:-

```
output "jenkins_public_ip" {
  description = "The public IP of the Jenkins server"
  value       = aws_instance.jenkins.public_ip
}

output "sonarqube_public_ip" {
  description = "The public IP of the SonarQube server"
  value       = aws_instance.sonarqube.public_ip
}
```

c. provider.tf:-

```
provider "aws" {
  region = var.aws_region
  access_key = var.aws_access_key
  secret_key = var.aws_secret_key
}
```

d. security_groups.tf:-

```
resource "aws_security_group" "jenkins_sg" {
  name        = "jenkins_security_group"
  description = "Allow SSH and HTTP access for Jenkins"

  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

```
ingress {
  from_port = 8080
  to_port   = 8080
  protocol  = "tcp"
  cidr_blocks = ["0.0.0.0/0"]
}

egress {
  from_port = 0
  to_port   = 0
  protocol  = "-1"
  cidr_blocks = ["0.0.0.0/0"]
}
}

resource "aws_security_group" "sonarqube_sg" {
  name      = "sonarqube_security_group"
  description = "Allow SSH and HTTP access for SonarQube"

  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port = 9000
    to_port   = 9000
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

e. variables.tf:-

```
variable "aws_region" {  
  description = "The AWS region to deploy resources in"  
  default     = "us-east-1"  
}  
variable "aws_access_key" {  
  description = "Your AWS access key"  
}  
variable "aws_secret_key" {  
  description = "Your AWS secret key"  
}  
variable "instance_type" {  
  description = "The EC2 instance type"  
  default     = "t2.micro"  
}  
variable "key_pair_name" {  
  description = "The name of the key pair for SSH access"  
}  
variable "jenkins_ami_id" {  
  description = "The AMI ID for the Jenkins server"  
}  
variable "sonarqube_ami_id" {  
  description = "The AMI ID for the SonarQube server"  
}
```

f. terraform.tfvars:-

```
aws_access_key = "YOUR_ACCESS_KEY_HERE"  
aws_secret_key = "YOUR_SECRET_KEY_HERE"  
aws_region     = "us-east-1"  
instance_type  = "t2.micro"  
key_pair_name  = "my-ssh-key"  
jenkins_ami_id = "ami-0abcdef1234567890"  
sonarqube_ami_id = "ami-0abcdef1234567890"
```

In the above file, replace the placeholder text with your values as such:-

```
terraform.tfvars > ...
1  aws_access_key = "AKIAQFC27GEZP2JF5H4K"
2  aws_secret_key = "EtMHs54TscjmSd1ykv2iveQqcwW7vFruRHM2iAY6"
3  aws_region = "us-east-1"
4  instance_type = "t2.micro"
5  key_pair_name = "my-ssh-key"
6  jenkins_ami_id = "ami-06b21ccaeff8cd686"
7  sonarqube_ami_id = "ami-06b21ccaeff8cd686"
8  |
```

Step 2: Open a terminal inside of the folder that was created for your project and run the following commands in the terminal:-

1. terraform init

```
PS C:\Users\anish\OneDrive\Desktop\AdvDevOps Case Study> terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/aws...
- Installing hashicorp/aws v5.72.1...
- Installed hashicorp/aws v5.72.1 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

2. terraform plan

```
PS C:\Users\anish\OneDrive\Desktop\AdvDevOps Case Study> terraform plan
```

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
# aws_instance.jenkins will be created
+ resource "aws_instance" "jenkins" {
  + ami                = "ami-06b21ccaeff8cd686"
  + arn                = (known after apply)
  + associate_public_ip_address = (known after apply)
  + availability_zone   = (known after apply)
  + cpu_core_count      = (known after apply)
  + cpu_threads_per_core = (known after apply)
  + disable_api_stop    = (known after apply)
  + disable_api_termination = (known after apply)
  + ebs_optimized       = (known after apply)
  + get_password_data   = false
  + host_id             = (known after apply)
  + host_resource_group_arn = (known after apply)
  + iam_instance_profile = (known after apply)
  + id                 = (known after apply)
  + instance_initiated_shutdown_behavior = (known after apply)
  + instance_lifecycle  = (known after apply)
  + instance_state      = (known after apply)
  + instance_type       = "t2.micro"
  + ipv6_address_count  = (known after apply)
```

```
    ]
    + from_port        = 9000
    + ipv6_cidr_blocks = []
    + prefix_list_ids  = []
    + protocol         = "tcp"
    + security_groups  = []
    + self             = false
    + to_port          = 9000
    # (1 unchanged attribute hidden)
  },
]
+ name                = "sonarqube_security_group"
+ name_prefix         = (known after apply)
+ owner_id            = (known after apply)
+ revoke_rules_on_delete = false
+ tags_all            = (known after apply)
+ vpc_id              = (known after apply)
}
```

Plan: 4 to add, 0 to change, 0 to destroy.

Changes to Outputs:

```
+ jenkins_public_ip = (known after apply)
+ sonarqube_public_ip = (known after apply)
```

Note: You didn't use the `-out` option to save this plan, so Terraform can't guarantee to take exactly these actions if you run `"terraform apply"` now.

3. terraform apply:-

```
PS C:\Users\anish\OneDrive\Desktop\AdvDevOps Case Study> terraform apply
```

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
# aws_instance.jenkins will be created
+ resource "aws_instance" "jenkins" {
  + ami               = "ami-06b21ccaeff8cd686"
  + arn               = (known after apply)
  + associate_public_ip_address = (known after apply)
  + availability_zone  = (known after apply)
  + cpu_core_count    = (known after apply)
  + cpu_threads_per_core = (known after apply)
  + disable_api_stop   = (known after apply)
  + disable_api_termination = (known after apply)
  + ebs_optimized      = (known after apply)
  + get_password_data   = false
  + host_id            = (known after apply)
  + host_resource_group_arn = (known after apply)
  + iam_instance_profile = (known after apply)
  + id                = (known after apply)
  + instance_initiated_shutdown_behavior = (known after apply)
  + instance_lifecycle = (known after apply)
  + instance_state     = (known after apply)
  + instance_type      = "t2.micro"
  + ipv6_address_count  = (known after apply)
  + ipv6_addresses     = (known after apply)
```

Plan: 4 to add, 0 to change, 0 to destroy.

Changes to Outputs:

```
+ jenkins_public_ip = (known after apply)
+ sonarqube_public_ip = (known after apply)
```

Do you want to perform these actions?

Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes

```
aws_security_group.jenkins_sg: Creating...
aws_security_group.sonarqube_sg: Creating...
aws_security_group.sonarqube_sg: Creation complete after 6s [id=sg-03a2c0a7016aae07e]
aws_security_group.jenkins_sg: Creation complete after 6s [id=sg-0630ec4a658d7b152]
aws_instance.jenkins: Creating...
aws_instance.sonarqube: Creating...
aws_instance.jenkins: Still creating... [10s elapsed]
aws_instance.sonarqube: Still creating... [10s elapsed]
aws_instance.jenkins: Creation complete after 15s [id=i-03d81881d1246aa2f]
aws_instance.sonarqube: Creation complete after 15s [id=i-0fd5644579ef07b0e]
```

Apply complete! Resources: 4 added, 0 changed, 0 destroyed.

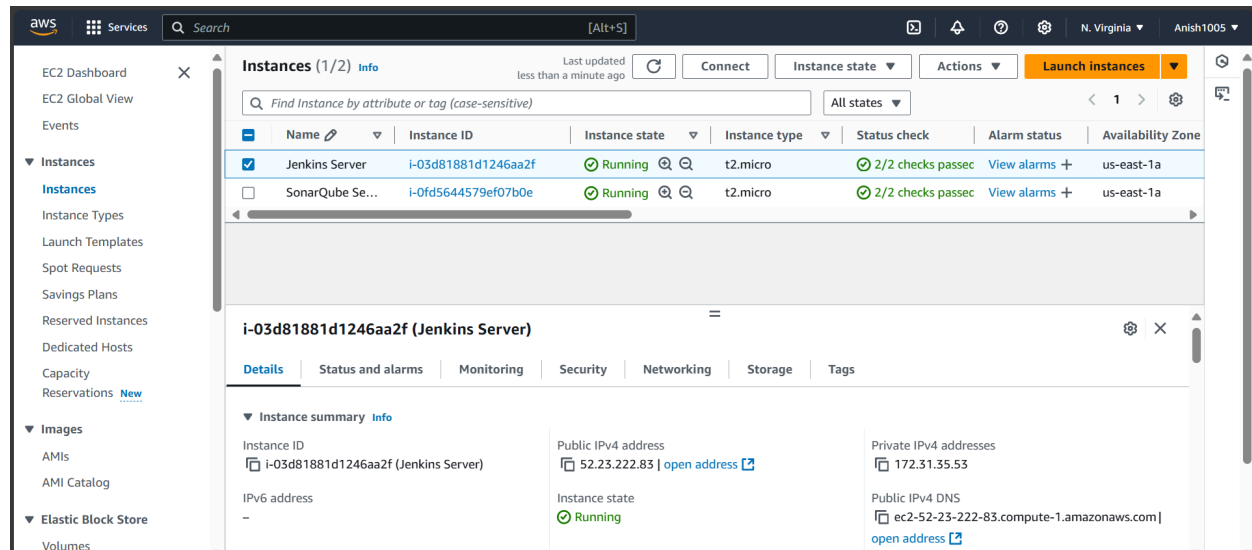
Outputs:

```
jenkins_public_ip = "52.23.222.83"
sonarqube_public_ip = "54.210.254.92"
```

```
PS C:\Users\anish\OneDrive\Desktop\AdvDevOps Case Study> |
```

‘terraform init’ initializes the working directory, ‘terraform plan’ previews changes without applying them, and ‘terraform apply’ executes the planned changes to create or modify infrastructure.

As a result of running the above, two instances (one for Jenkins server and the other for the SonarQube server) are created on our AWS account.



Step 3: Connect both the instances to your local terminal using SSH.

Step 4: Go onto your AWS console, change the instance state of both the instances to 'Stopped' and change the instance type of 'Jenkins Server' from t2.micro to t3.medium and change the instance type of 'SonarQube Server' from t2.micro to t3.small. This must be done because both Jenkins and SonarQube require a lot of space (in GBs) and t2.micro doesn't have enough memory space to accommodate either Jenkins or SonarQube. Once this is done, change the instance type of both the instances to 'Running'.

Make sure to stop your instances when they aren't in use because t3.micro is significantly more expensive than t2.micro and keeping them running can add to your bill significantly.

Step 5: Open the terminal for your Jenkins server and run the following commands to install Jenkins onto your instance:-

```
sudo yum update -y
sudo wget -O /etc/yum.repos.d/jenkins.repo https://pkg.jenkins.io/redhat-stable/jenkins.repo
sudo rpm --import https://pkg.jenkins.io/redhat-stable/jenkins.io.key
sudo yum install java-17-amazon-corretto
sudo yum install jenkins -y --nogpgcheck
sudo systemctl start jenkins
sudo systemctl enable jenkins
```

Running the above commands should install Jenkins onto your instance. Now, to open Jenkins, paste the following into your browser:-

`http://<public_ip_address_of_your_Jenkins_instance>:8080`

The Jenkins Sign Up page should appear. Complete the configuration:

Getting Started

Create First Admin User

Username

Password

Confirm password

Jenkins 2.462.3


[Skip and continue as admin](#)[Save and Continue](#)

Getting Started

Confirm password

Full name

E-mail address

 Invalid e-mail address

Jenkins 2.462.3

[Skip and continue as admin](#)[Save and Continue](#)

Getting Started

Instance Configuration

Jenkins URL:

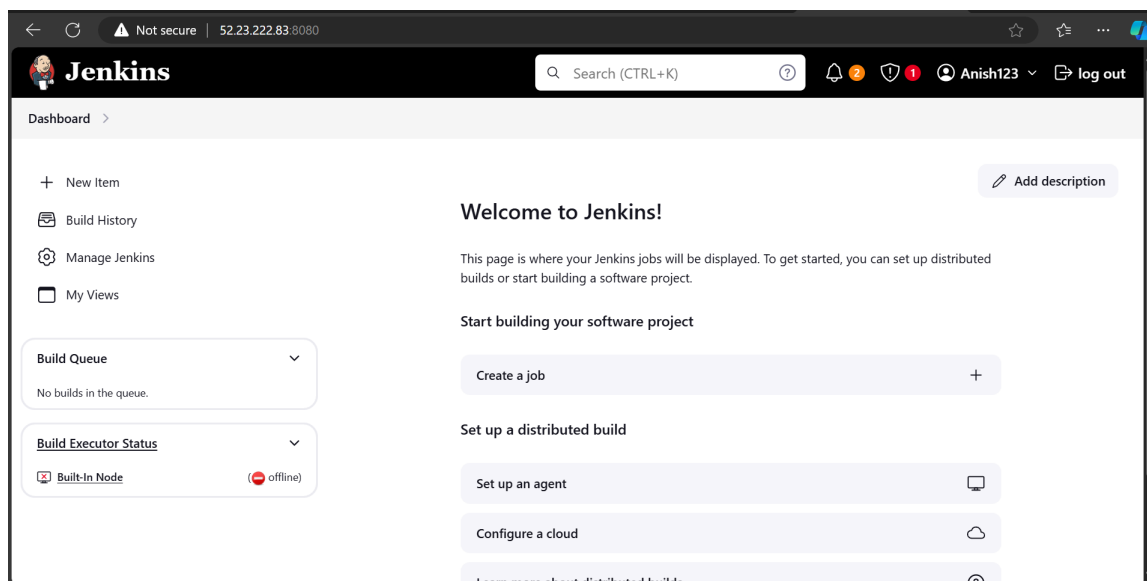
The Jenkins URL is used to provide the root URL for absolute links to various Jenkins resources. That means this value is required for proper operation of many Jenkins features including email notifications, PR status updates, and the `BUILD_URL` environment variable provided to build steps.

The proposed default value shown is **not saved yet** and is generated from the current request, if possible. The best practice is to set this value to the URL that users are expected to use. This will avoid confusion when sharing or viewing links.

Jenkins 2.462.3

Not now

Save and Finish



The Jenkins dashboard should open.

Step 6: Open the terminal for your SonarQube server and run the following commands to install SonarQube onto your instance:-

```
sudo wget https://binaries.sonarsource.com/Distribution/sonarqube/sonarqube-9.9.2.77730.zip  
-P /opt
```

```
sudo unzip /opt/sonarqube-9.9.2.77730.zip -d /opt
```

```
sudo mv /opt/sonarqube-9.9.2.77730 /opt/sonarqube
```

```
sudo chown -R sonar:sonar /opt/sonarqube
```

```
cd /opt/sonarqube/bin/linux-x86-64
```

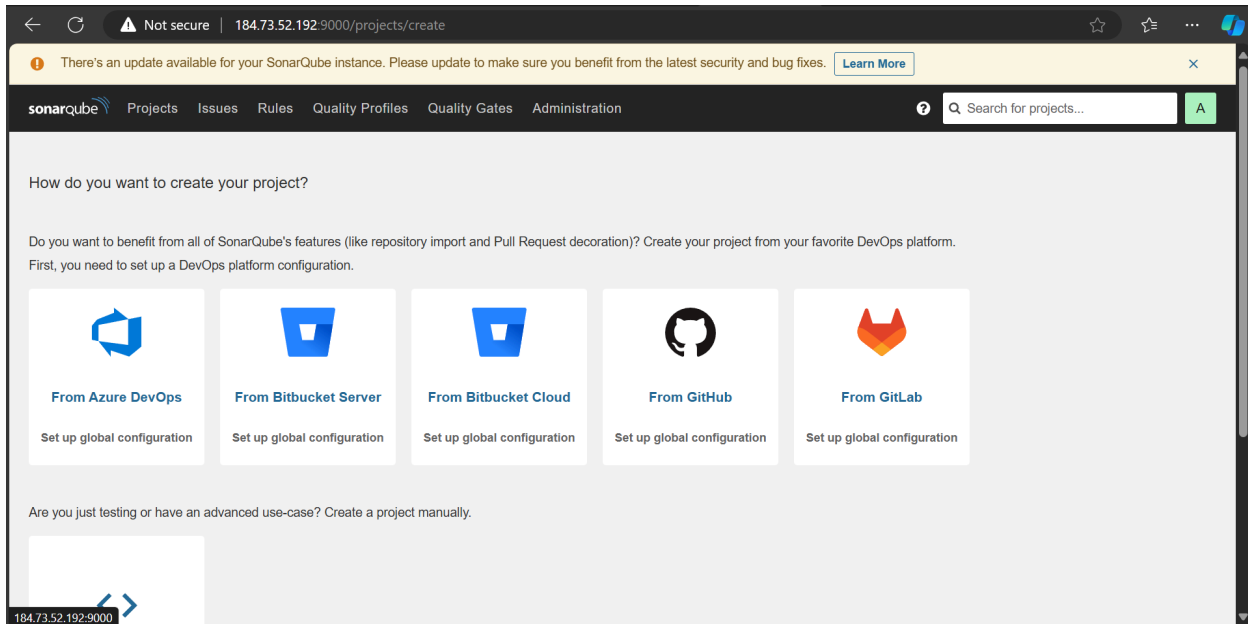
```
./sonar.sh start
```

Running the above commands should install SonarQube onto your instance. Now, to open SonarQube, paste the following into your browser:-

`http://<public_ip_address_of_your_SonarQube_instance>:9000`

The SonarQube Sign Up page should appear. The default credentials are admin/admin. Then, change the password to the password of your choice and complete the configuration.

After completing all of the above, the SonarQube dashboard should appear:-




Step 7: Create a SonarQube project by clicking on 'Manually' on the dashboard and giving it a name and completing its configuration:-


Create a project

All fields marked with * are required

Project display name *

 
Up to 255 characters. Some scanners might override the value you provide.

Project key *

 
The project key is a unique identifier for your project. It may contain up to 400 characters. Allowed characters are alphanumeric, '-' (dash), '_' (underscore), '.' (period) and ':' (colon), with at least one non-digit.

Main branch name *

The name of your project's default branch [Learn More](#)

[Set Up](#)

Step 8: Install SonarScanner CLI onto your system.

Step 9: Go onto the terminal of your Jenkins server and run the following commands:-

```
scp -i "<location_of_your_pem_file>" "location_of_sonar-scanner_cli"
ec2-user@<ip_address_of_SonarQube_server>:/home/ec2-user/
unzip sonar-scanner-cli-6.2.0.4584-windows-x64.zip
```

```
[ec2-user@ip-172-31-35-53 ~]$ unzip sonar-scanner-cli-6.2.1.4610-linux-x64.zip
Archive:  sonar-scanner-cli-6.2.1.4610-linux-x64.zip
  creating: sonar-scanner-6.2.1.4610-linux-x64/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/conf/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/conf/security/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/conf/security/policy/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/conf/security/policy/unlimited/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/conf/security/policy/limited/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/conf/sdp/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/conf/management/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/lib/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/lib/server/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/lib/security/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/lib/jfr/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/legal/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/legal/jdk.incubator.vector/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/legal/jdk.charsets/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.scripting/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/legal/jdk.security.auth/
  creating: sonar-scanner-6.2.1.4610-linux-x64/jre/legal/jdk.xml.dom/
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.datatransfer/LICENSE -> ../java.base/LICENSE
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.rmi/ASSEMBLY_EXCEPTION -> ../java.base/ASSEMBLY_EXCEPTION
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.rmi/ADDITIONAL_LICENSE_INFO -> ../java.base/ADDITIONAL_LICENSE_INFO
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.rmi/LICENSE -> ../java.base/LICENSE
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.logging/ASSEMBLY_EXCEPTION -> ../java.base/ASSEMBLY_EXCEPTION
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.logging/ADDITIONAL_LICENSE_INFO -> ../java.base/ADDITIONAL_LICENSE_INFO
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.logging/LICENSE -> ../java.base/LICENSE
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.management.rmi/ASSEMBLY_EXCEPTION -> ../java.base/ASSEMBLY_EXCEPTION
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.management.rmi/ADDITIONAL_LICENSE_INFO -> ../java.base/ADDITIONAL_LICENSE_INFO
  sonar-scanner-6.2.1.4610-linux-x64/jre/legal/java.management.rmi/LICENSE -> ../java.base/LICENSE
[ec2-user@ip-172-31-35-53 ~]$
```

```
sudo mv sonar-scanner-6.2.1.4610-linux-x64 /opt/sonar-scanner
echo 'export PATH=$PATH:/opt/sonar-scanner/bin' >> ~/.bashrc
source ~/.bashrc
sonar-scanner -v
```

```
[ec2-user@ip-172-31-35-53 ~]$ sudo mv sonar-scanner-6.2.1.4610-linux-x64 /opt/sonar-scanner
[ec2-user@ip-172-31-35-53 ~]$ echo 'export PATH=$PATH:/opt/sonar-scanner/bin' >> ~/.bashrc
source ~/.bashrc
sonar-scanner -v
05:07:33.776 INFO Scanner configuration file: /opt/sonar-scanner/conf/sonar-scanner.properties
05:07:33.782 INFO Project root configuration file: NONE
05:07:33.805 INFO SonarScanner CLI 6.2.1.4610
05:07:33.808 INFO Java 17.0.12 Eclipse Adoptium (64-bit)
05:07:33.809 INFO Linux 6.1.112-122.189.amzn2023.x86_64 amd64
[ec2-user@ip-172-31-35-53 ~]$
```

nano /opt/sonar-scanner/sonar-scanner-6.2.1.4610-linux-x64/conf/sonar-scanner.properties

When the sonar.properties file opens, add the following to the file:-

sonar.host.url=http://<your-sonarqube-ip>:9000

Step 10: Go to Manage Jenkins > Plugins from your Jenkins dashboard and install the following plugins:-

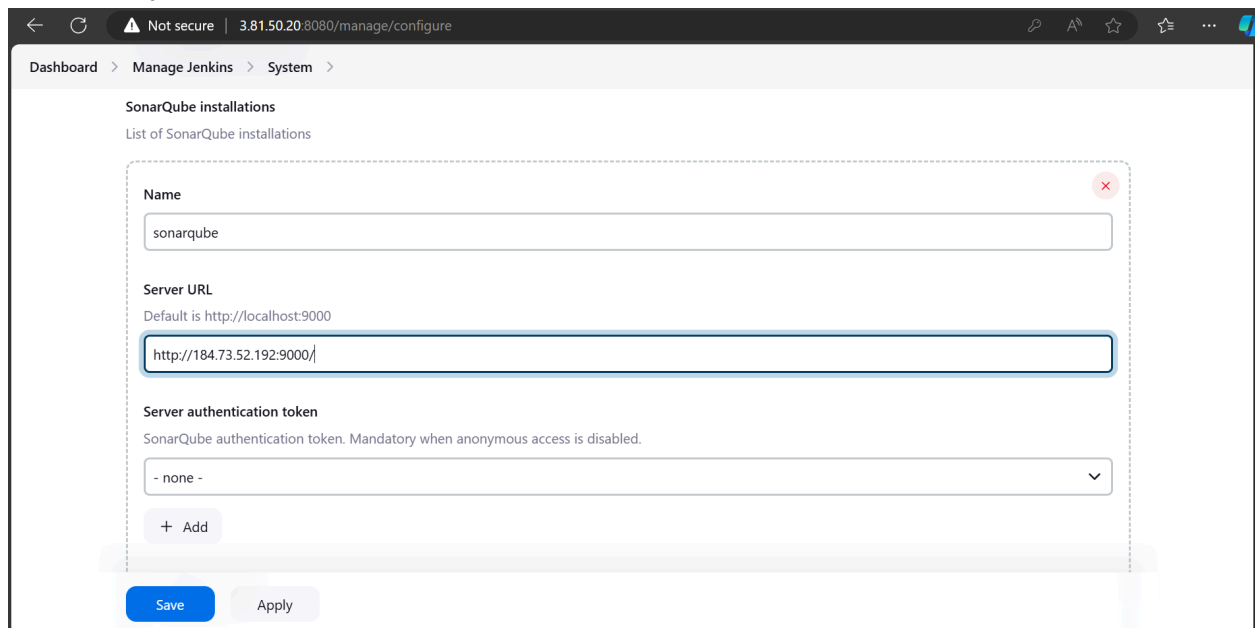
Pipeline: Stage View

Pipeline: Rest API

SonarQube Scanner

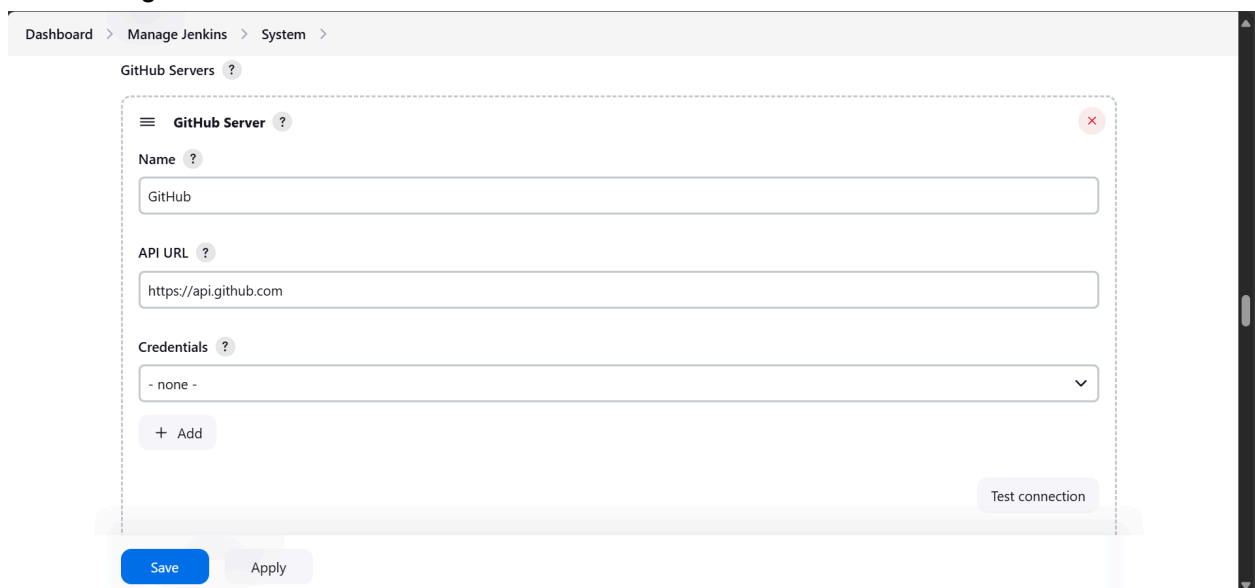
Git server

Step 11: Go to Manage Jenkins > System from your Jenkins dashboard, scroll down to 'SonarQube installations' and click on 'Add SonarQube'. Then give it a name and add the IP address of your SonarQube.



The screenshot shows the Jenkins 'System' configuration page, specifically the 'SonarQube installations' section. The breadcrumb trail is 'Dashboard > Manage Jenkins > System > SonarQube installations'. Below the section title, it says 'List of SonarQube installations'. A dashed box contains the configuration form for a new installation. The 'Name' field is filled with 'sonarqube'. The 'Server URL' field has a default value of 'http://localhost:9000' and is currently filled with 'http://184.73.52.192:9000/'. The 'Server authentication token' dropdown is set to '- none -'. At the bottom of the dashed box is a '+ Add' button. Below the dashed box are 'Save' and 'Apply' buttons.

Then, navigate to 'GitHub Servers' and write the name as 'GitHub'.



The screenshot shows the Jenkins 'System' configuration page, specifically the 'GitHub Servers' section. The breadcrumb trail is 'Dashboard > Manage Jenkins > System > GitHub Servers'. Below the section title, it says 'GitHub Servers ?'. A dashed box contains the configuration form for a new server. The 'Name' field is filled with 'GitHub'. The 'API URL' field is filled with 'https://api.github.com'. The 'Credentials' dropdown is set to '- none -'. At the bottom of the dashed box is a '+ Add' button. To the right of the dashed box is a 'Test connection' button. Below the dashed box are 'Save' and 'Apply' buttons.

Step 12: Go to Manage Jenkins > Tools from your Jenkins dashboard and scroll down to 'SonarQube Scanner installations' and click on 'Add SonarQube Scanner'. Then, give it a name and choose the default version and click on 'Save'.

Dashboard > Manage Jenkins > Tools

SonarQube Scanner installations

Add SonarQube Scanner

SonarQube Scanner

Name

SonarScanner

☒ Install automatically ?

Install from Maven Central

Version

SonarQube Scanner 6.2.1.4610

Add Installer ▾

Save Apply

Step 13: On your Jenkins dashboard, click on 'New Item'. Then, give your project a name and click on 'Pipeline' as Item type.

Jenkins

Search (CTRL+K)

Dashboard > All > New Item

New Item

Enter an item name

CaseStudy29

Select an item type

Freestyle project
Classic, general-purpose job type that checks out from up to one SCM, executes build steps serially, followed by post-build steps like archiving artifacts and sending email notifications.

Pipeline
Orchestrates long-running activities that can span multiple build agents. Suitable for building pipelines (formerly known as workflows) and/or organizing complex activities that do not easily fit in free-style job type.

Multi-configuration project

OK

Then, in the Git section, add the link to the repository which contains the python project that you are going to test SonarQube on.

Next, scroll down to the 'Pipeline script' section and enter the following:-

```
node {
```

```
    stage('Cloning the GitHub Repo') {
        git branch: 'main', url: '<link_to_repository_of_python_project>'
    }
}
```

```
}

stage('SonarQube Analysis') {
    withSonarQubeEnv('sonarqube') {
        sh """
        <path_to_your_sonar-scanner>\
        -D sonar.login=admin \
        -D sonar.password=<your_password> \
        -D sonar.projectKey=<your_project_key> \
        -D sonar.exclusions=vendor/*,resources/,.java \
        -D sonar.host.url=<url_of_your_sonarqube>
        """
    }
}
}
```

The screenshot shows the SonarQube web interface for configuring a pipeline. The breadcrumb navigation at the top reads "Dashboard > CaseStudy29 > Configuration". On the left, under the "Configure" section, there are three tabs: "General" (selected), "Advanced Project Options", and "Pipeline". The main area is titled "Pipeline" and shows the "Definition" dropdown set to "Pipeline script". Below this is a "Script" editor with a line number indicator on the left (1-17) and a "try sample Pipeline..." button on the right. The script content is as follows:

```
1 node {
2   stage('Cloning the GitHub Repo') {
3     git branch: 'main', url: 'https://github.com/Rakshit5467/PulseAnalytics.git'
4   }
5
6   stage('SonarQube Analysis') {
7     withSonarQubeEnv('sonarqube') {
8       sh """
9         /opt/sonar-scanner/bin/sonar-scanner \
10        -D sonar.login=admin \
11        -D sonar.password=ANISH2004 \
12        -D sonar.projectKey=CaseStudyPython \
13        -D sonar.exclusions=vendor/*,resources/,.java \
14        -D sonar.host.url=http://184.73.52.192:9000/
15        """
16     }
17 }
```

At the bottom of the script editor, there is a checkbox labeled "Use Groovy Sandbox" which is checked. Below the editor are two buttons: "Save" (in blue) and "Apply" (in grey).

Then, click on Save.

Step 14: Go to the Jenkins dashboard and click on 'Build now' from the left sidebar.

Jenkins Search (CTRL+K) ? 1 Anish123 log out

Dashboard > CaseStudy29 >

CaseStudy29 Add description

Stage View

	Cloning the GitHub Repo	SonarQube Analysis
Average stage times: (Average full run time: ~38s)	4s	31s
#1 Oct 20 11:12 No Changes	4s	31s

Permalinks

Dashboard > CaseStudy29 >

- Delete Pipeline
- Full Stage View
- GitHub
- Stages
- Rename
- Pipeline Syntax

Build History trend

Filter...

#1 Oct 20, 2024, 5:42 AM

Atom feed for all Atom feed for failures

This builds your project.

You can also check your console output:-



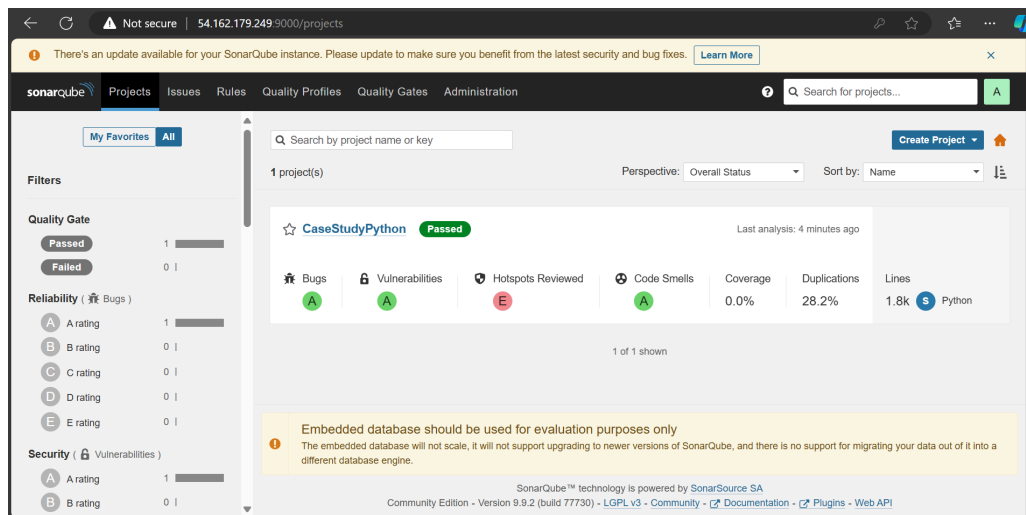
Console Output

[Download](#)[Copy](#)[View as plain text](#)

```
Started by user Anish123
[Pipeline] Start of Pipeline (hide)
[Pipeline] node
Running on Jenkins in /var/lib/jenkins/workspace/CaseStudy29
[Pipeline] {
[Pipeline] stage
[Pipeline] { (Cloning the GitHub Repo)
[Pipeline] git
The recommended git tool is: NONE
No credentials specified
Cloning the remote Git repository
Cloning repository https://github.com/Rakshit5467/PulseAnalytics.git
> /usr/bin/git init /var/lib/jenkins/workspace/CaseStudy29 # timeout=10
Fetching upstream changes from https://github.com/Rakshit5467/PulseAnalytics.git
> /usr/bin/git --version # timeout=10
> git --version # 'git version 2.40.1'
> /usr/bin/git fetch --tags --force --progress -- https://github.com/Rakshit5467/PulseAnalytics.git
```

```
05:43:00.771 INFO SCM Publisher 20/20 source files have been analyzed (done) | time=803ms
05:43:00.788 INFO CPD Executor Calculating CPD for 19 files
05:43:00.881 INFO CPD Executor CPD calculation finished (done) | time=93ms
05:43:01.033 INFO Analysis report generated in 142ms, dir size=413.4 kB
05:43:01.145 INFO Analysis report compressed in 109ms, zip size=198.9 kB
05:43:01.402 INFO Analysis report uploaded in 256ms
05:43:01.404 INFO ANALYSIS SUCCESSFUL, you can find the results at: http://184.73.52.192:9000/dashboard?id=CaseStudyPython
05:43:01.404 INFO Note that you will be able to access the updated dashboard once the server has processed the submitted analysis report
05:43:01.404 INFO More about the report processing at http://184.73.52.192:9000/api/ce/task?id=AZKoc0wD-fZOm3C89gfI
05:43:01.466 INFO Analysis total time: 20.739 s
05:43:01.488 INFO EXECUTION SUCCESS
05:43:01.490 INFO Total time: 30.191s
[Pipeline] }
[Pipeline] // withSonarQubeEnv
[Pipeline] }
[Pipeline] // stage
[Pipeline] }
[Pipeline] // node
[Pipeline] End of Pipeline
Finished: SUCCESS
```

Step 15: You can also go onto your SonarQube dashboard to check the output on your project.



Conclusion: This case study demonstrates how integrating Terraform, Jenkins, and SonarQube can effectively automate infrastructure setup and continuous code analysis. By using Terraform's Infrastructure-as-Code approach, the process of deploying Jenkins and SonarQube on AWS becomes scalable, consistent, and easy to manage, reducing manual effort. Jenkins automates tasks like building, testing, and deploying code, while SonarQube scans for issues such as bugs, code smells, and vulnerabilities, providing developers with early feedback. Together, these tools streamline the CI/CD pipeline, improving code quality and reducing the time spent on fixing problems.